Part 3: Relational Data, Preparing Data and Linear Models

(a) Relational Data with dplyr

Relational Data with dplyr

- Typically, data analysis involves many tables of data that must be combined to answer questions
- Collectively, multiple tables of data are called relational data
- Relations are always defined between a pair of tables
- See tibbles **x** and **y**

```
## # A tibble: 3 x 2
##     key val_x
##   <dbl> <chr>
## 1     1 x1
## 2     2 x2
## 3     3 x3
```

```
## # A tibble: 3 x 2
##     key val_y
##   <dbl> <chr>
## 1     1 y1
## 2     2 y2
## 3     4 y3
```

Keys

- The variables used to connect each pair of tables are called keys
- A key is a variable (or set of variables) that uniquely identifies an observation
- There are two types of keys:
  - A primary key uniquely identifies an observation in its own table
  - A foreign key uniquely identifies an observation in another table.

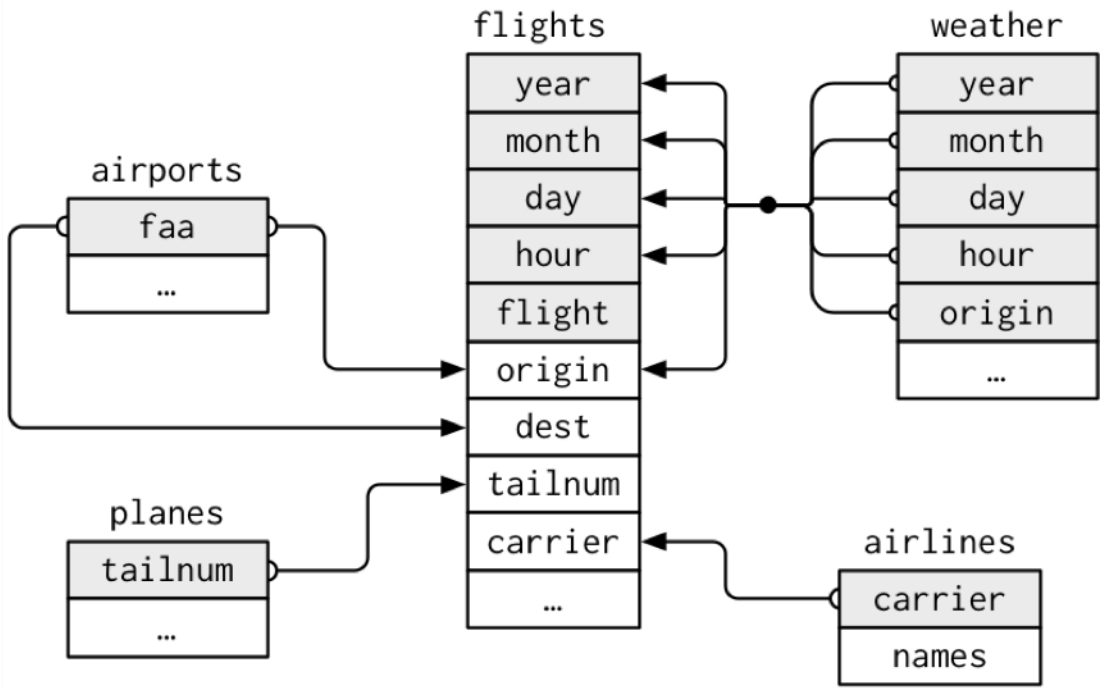

Figure 1: Keys and Relations in flights tables

Mutating Joins

- Allows you to combine variables from two tables
- First matches observations by their keys, and then copies across variables from one table to another
- Similar to mutate(), the join functions add variables to the right
- Types
  - Inner Join
  - Left Join
  - Right Join
  - Full Join

Inner Joins

- Matches pairs of observations when their keys are equal
- Unmatched rows are not included in the result

```
inner_join(x,y)
```

```
## Joining, by = "key"
```

```
## # A tibble: 2 x 3
##     key val_x val_y
##   <dbl> <chr> <chr>
## 1     1 x1    y1
## 2     2 x2    y2
```

x

key	val_x
1	x1
2	x2
3	x3

y

key	val_y
1	y1
2	y2
4	y3

Figure 2: Tables x and y

Left Join

A left join keeps all observations in x

```
left_join(x,y)
```

```
## Joining, by = "key"
```

```
## # A tibble: 3 x 3
##     key val_x val_y
##   <dbl> <chr> <chr>
## 1     1 x1    y1
## 2     2 x2    y2
## 3     3 x3    <NA>
```

x

key	val_x
1	x1
2	x2
3	x3

y

key	val_y
1	y1
2	y2
4	y3

Figure 3: Tables x and y

Right Join

A right join keeps all observations in y

```
right_join(x,y)
```

```
## Joining, by = "key"
```

```
## # A tibble: 3 x 3
##     key val_x val_y
##   <dbl> <chr> <chr>
## 1     1 x1    y1
## 2     2 x2    y2
## 3     4 <NA>  y3
```

x

key	val_x
1	x1
2	x2
3	x3

y

key	val_y
1	y1
2	y2
4	y3

Figure 4: Tables x and y

Full Join

A full join keeps all observations in x and y

```
full_join(x,y)
```

```
## Joining, by = "key"
```

```
## # A tibble: 4 x 3
##     key val_x val_y
##   <dbl> <chr> <chr>
## 1     1 x1    y1
## 2     2 x2    y2
## 3     3 x3    <NA>
## 4     4 <NA>  y3
```

x

key	val_x
1	x1
2	x2
3	x3

y

key	val_y
1	y1
2	y2
4	y3

Figure 5: Tables x and y

Filtering Joins

Match observations in the same way as mutating joins, but affect the observations, not the variables. Two types:

- semi_join(x,y) keeps all observations in x that have a match in y
- anti_join(x,y), drops all observations in x that have a match in y.

Semi Joins

Keeps all observations in x that have a match in y

```
semi_join(x,y)
```

```
## Joining, by = "key"
```

```
## # A tibble: 2 x 2
##     key val_x
##   <dbl> <chr>
## 1     1 x1
## 2     2 x2
```

x

key	val_x
1	x1
2	x2
3	x3

y

key	val_y
1	y1
2	y2
4	y3

Figure 6: Tables x and y

Anti Joins

Drops all observations in x that have a match in y.

```
anti_join(x,y)
```

```
## Joining, by = "key"
```

```
## # A tibble: 1 x 2
##     key val_x
##   <dbl> <chr>
## 1     3 x3
```

x

key	val_x
1	x1
2	x2
3	x3

y

key	val_y
1	y1
2	y2
4	y3

Figure 7: Tables x and y

Challenge 3.1

- Filter out incomplete flights from the dataset
- Join the flights data to the weather data
- Filter out missing temperature values
- Plot the relationship between temperatures and departure delays, facet by origin and colour by wind_speed
- Use a sample of 10000 for the plot, with seed 99.

Test Slide with Plot