

## 6. Exploratory Data Analysis - dplyr

Data Science for OR - J. Duggan

### Overview

- Visualisation is an important tool for insight generation, but it's rare that you get the data in exactly the right form you need" (Wickham and Grolemund 2017)
  - Create new variables
  - Create summaries
  - Order data
- dplyr package is designed for data transformation

- All verbs (functions) work similarly
- The first argument is a data frame/tibble
- The subsequent arguments decide what to do with the data frame
- The result is a data frame (supports chaining of steps)

Function	Purpose
<b>filter()</b>	Pick observations by their values
<b>arrange()</b>	Reorder the rows
<b>select()</b>	Pick variables by their names
<b>mutate()</b>	Create new variables with functions of existing variables
<b>summarise()</b>	Collapse many values down to a single summary

## Sample Data set ggplot2::mpg

```
## Observations: 234
## Variables: 11
## $ manufacturer <chr> "audi", "audi", "audi", "audi", "audi"
## $ model        <chr> "a4", "a4", "a4", "a4", "a4", "a4", "a4"
## $ displ        <dbl> 1.8, 1.8, 2.0, 2.0, 2.8, 2.8, 3.1, 1.8
## $ year         <int> 1999, 1999, 2008, 2008, 1999, 1999, 2008
## $ cyl          <int> 4, 4, 4, 4, 6, 6, 6, 4, 4, 4, 4, 6, 6
## $ trans        <chr> "auto(l5)", "manual(m5)", "manual(m6)"
## $ drv          <chr> "f", "f", "f", "f", "f", "f", "f", "4"
## $ cty          <int> 18, 21, 20, 21, 16, 18, 18, 18, 16, 20
## $ hwy          <int> 29, 29, 31, 30, 26, 26, 27, 26, 25, 28
## $ fl           <chr> "p", "p", "p", "p", "p", "p", "p", "p"
## $ class        <chr> "compact", "compact", "compact", "compact"
```

## (1) filter()

- Subset observations based on their values.
- First argument the name of the data frame
- Subsequent arguments are expressions that filter the data frame
- Only includes rows that have no missing values

```
filter(mpg,manufacturer=="audi",year==1999,model=="a4")
```

```
## # A tibble: 4 x 11
##   manufacturer model displ  year   cyl trans  drv      cty
##   <chr>          <chr> <dbl> <int> <int> <chr> <chr> <int> <chr>
## 1 audi          a4      1.8  1999     4 auto(~ f      18
## 2 audi          a4      1.8  1999     4 manua~ f      21
## 3 audi          a4      2.8  1999     6 auto(~ f      16
## 4 audi          a4      2.8  1999     6 manua~ f      18
```

## Cars with highest mpg, lowest mpg?

```
filter(mpg,hwy==max(hwy))
```

```
## # A tibble: 2 x 11
##   manufacturer model displ  year   cyl trans  drv      cty
##   <chr>          <chr> <dbl> <int> <int> <chr> <chr> <int> <chr>
## 1 volkswagen    jetta    1.9  1999     4 manu~ f      33
## 2 volkswagen    new b~    1.9  1999     4 manu~ f      35
```

```
filter(mpg,hwy==min(hwy))
```

```
## # A tibble: 5 x 11
##   manufacturer model displ  year   cyl trans  drv      cty
##   <chr>          <chr> <dbl> <int> <int> <chr> <chr> <int> <chr>
## 1 dodge         dakot~    4.7  2008     8 auto~ 4      9
## 2 dodge         duran~    4.7  2008     8 auto~ 4      9
## 3 dodge         ram 1~    4.7  2008     8 auto~ 4      9
## 4 dodge         ram 1~    4.7  2008     8 manu~ 4      9
```

## Challenge 2.1

- List the cars with an average city mpg greater than the median.
- Show the cars with the maximum displacement

### (2) arrange()

- Changes the order of rows.
- Takes a data frame and a set of column names to order by

```
arrange(mpg, displ)
```

```
## # A tibble: 234 x 11
##   manufacturer model displ year   cyl trans drv      cty
##   <chr>          <chr> <dbl> <int> <int> <chr> <chr> <int> <
## 1 honda         civic   1.6  1999     4 manu~ f      28
## 2 honda         civic   1.6  1999     4 auto~ f      24
## 3 honda         civic   1.6  1999     4 manu~ f      25
## 4 honda         civic   1.6  1999     4 manu~ f      23
## 5 honda         civic   1.6  1999     4 auto~ f      24
## 6 audi          a4      1.8  1999     4 auto~ f      18
## 7 audi          a4      1.8  1999     4 manu~ f      21
## 8 audi          a4 q~   1.8  1999     4 manu~ 4      18
## 9 audi          a4 q~   1.8  1999     4 auto~ 4      16
```

## Show in descending order

```
arrange(mpg,desc(displ))
```

```
## # A tibble: 234 x 11
##   manufacturer model displ year   cyl trans drv   cty
##   <chr>          <chr> <dbl> <int> <int> <chr> <chr> <int> <
## 1 chevrolet      corv~    7   2008     8 manu~ r    15
## 2 chevrolet      k150~   6.5  1999     8 auto~ 4    14
## 3 chevrolet      corv~   6.2  2008     8 manu~ r    16
## 4 chevrolet      corv~   6.2  2008     8 auto~ r    15
## 5 jeep           gran~   6.1  2008     8 auto~ 4    11
## 6 chevrolet      c150~    6   2008     8 auto~ r    12
## 7 dodge          dura~   5.9  1999     8 auto~ 4    11
## 8 dodge          ram ~   5.9  1999     8 auto~ 4    11
## 9 chevrolet      c150~   5.7  1999     8 auto~ r    13
## 10 chevrolet     corv~   5.7  1999     8 manu~ r    16
## # ... with 224 more rows
```

## Add an extra sort column

```
arrange(mpg,desc(year),desc(displ))
```

```
## # A tibble: 234 x 11
##   manufacturer model displ year   cyl trans drv   cty
##   <chr>          <chr> <dbl> <int> <int> <chr> <chr> <int> <
## 1 chevrolet      corv~    7   2008     8 manu~ r    15
## 2 chevrolet      corv~   6.2  2008     8 manu~ r    16
## 3 chevrolet      corv~   6.2  2008     8 auto~ r    15
## 4 jeep           gran~   6.1  2008     8 auto~ 4    11
## 5 chevrolet      c150~    6   2008     8 auto~ r    12
## 6 dodge          dura~   5.7  2008     8 auto~ 4    13
## 7 dodge          ram ~   5.7  2008     8 auto~ 4    13
## 8 jeep           gran~   5.7  2008     8 auto~ 4    13
## 9 toyota         land~   5.7  2008     8 auto~ 4    13
## 10 nissan         path~   5.6  2008     8 auto~ 4    12
## # ... with 224 more rows
```

## (3) select()

- It is not uncommon to get datasets with hundreds, or even thousands, of variables
- A challenge is to narrow down on the variables of you're interested in
- `select()` allows you to rapidly zoom in on a useful subset using operations based on the variable names

```
select(mpg,model,year,displ, cty, hwy)
```

```
## # A tibble: 234 x 5
##   model      year displ   cty   hwy
##   <chr>    <int> <dbl> <int> <int>
## 1 a4      1999   1.8    18    29
## 2 a4      1999   1.8    21    29
## 3 a4      2008    2     20    31
## 4 a4      2008    2     21    30
## 5 a4      1999   2.8    16    26
## 6 a4      1999   2.8    18    26
```

## Special Function with select

### Special functions

As well as using existing functions like `:` and `c`, there are a number of special functions that only work inside `select`

- `starts_with(x, ignore.case = TRUE)`: names starts with `x`
- `ends_with(x, ignore.case = TRUE)`: names ends in `x`
- `contains(x, ignore.case = TRUE)`: selects all variables whose name contains `x`
- `matches(x, ignore.case = TRUE)`: selects all variables whose name matches the regular expression `x`
- `num_range("x", 1:5, width = 2)`: selects all variables (numerically) from `x01` to `x05`.
- `one_of("x", "y", "z")`: selects variables provided in a character vector.
- `everything()`: selects all variables.

## (4) mutate()

- It is often useful to add new columns that are functions of existing columns
- mutate() always adds new columns at the end of your data set.

```
sml <- select(mpg,model,displ,cty)
sml <- mutate(sml,Category=ifelse(cty>mean(cty),"AboveAvr","BelowAvr")
sml
```

```
## # A tibble: 234 x 4
##   model      displ  cty Category
##   <chr>      <dbl> <int> <chr>
## 1 a4          1.8    18 AboveAvr
## 2 a4          1.8    21 AboveAvr
## 3 a4          2      20 AboveAvr
## 4 a4          2      21 AboveAvr
## 5 a4          2.8    16 BelowAvr
## 6 a4          2.8    18 AboveAvr
```

## Useful creation functions

- There are many functions for creating new variables that can be used with mutate()
- The key property is that the function must be vectorised:
  - It must take a vector of values as input, and,
  - Return a vector with the same number of values as output

Grouping	Examples
<b>Arithmetic Operators</b>	+, -, *, /, ^
<b>Modular Arithmetic</b>	%% - Integer division && - Remainder
<b>Logs</b>	log(), log2(), log10()
<b>Offsets</b>	lead() and lag() Find when values change x!=lag(x)
<b>Cumulative and rolling aggregates</b>	cumsum(), cumprod(), cummin(), cummax(), cummean()
<b>Logical comparisons</b>	<, <=, >, >=, !=
<b>Ranking</b>	min_rank()

## (5) summarise()

- The last key verb is summarise()
- It collapses a data frame into a single row
- Not very useful unless paired with group\_by()
- Very useful to combine with the pipe operator %>%
- The pipe %>% comes from the magrittr package (Stefan Milton Bache)
- Helps to write code that is easier to read and understand
  - $x \%>\% f(y)$  turns into  $f(x, y)$

```
mpg %>% select(model, displ, cty) %>% slice(1:2)
```

```
## # A tibble: 2 x 3
##   model displ  cty
##   <chr> <dbl> <int>
## 1 a4      1.8    18
## 2 a4      1.8    21
```

## The function group\_by()

- Most summary data operations are useful done on groups defined by variables in the the dataset.
- The group\_by function takes an existing tbl and converts it into a grouped tbl where operations can then performed “by group”.

```
gr <- group_by(mpg, year)
agg <- summarise(gr, AverageCty=mean(cty))
agg
```

```
## # A tibble: 2 x 2
##   year AverageCty
##   <int>      <dbl>
## 1  1999      17.0
## 2  2008      16.7
```




## Using %>%

```
mpg %>% group_by(manufacturer) %>%  
  summarise(AvrCty=mean(cty),N=n()) %>%  
  arrange(desc(AvrCty)) %>%  
  slice(1:5)
```

```
## # A tibble: 5 x 3  
##   manufacturer AvrCty     N  
##   <chr>        <dbl> <int>  
## 1 honda         24.4     9  
## 2 volkswagen    20.9    27  
## 3 subaru        19.3    14  
## 4 hyundai       18.6    14  
## 5 toyota        18.5    34
```


## Overall idea

Original data frame



group\_by()

Grouped data frame



summarise()



# Useful Summary Functions

Grouping	Examples
Measures of location	mean(), median()
Measures of spread	sd(), IQR(),mad()
Measures of rank	min(), quantile(), max()
Measures of position	first(), nth(), last()
Counts	n(), n_distinct()
Counts and proportions of logical values	sum(x>0) when used with numeric functions, (T,F) converted to (1,0)

## The package nycflights13

```
glimpse(nycflights13::flights)
```


```
## Observations: 336,776
## Variables: 19
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013,
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
## $ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 5
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 6
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2,
## $ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913,
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854,
## $ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -
## $ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA",
## $ flight    <int> 1545, 1714, 1141, 725, 461, 1696, 50
## $ tailnum   <chr> "N14228", "N24211", "N619AA", "N804"
```

## Challenge 2.2 | nycflights13::flights

Generate the following graph. Use the variable **dep\_delay**. The variable **origin** indicates the departure airport.

```
unique(nycflights13::flights$origin)
```

```
## [1] "EWR" "LGA" "JFK"
```



## Summary

- dplyr - a grammar of data manipulation
- Five verbs
  - **filter()**
  - **arrange()**
  - **select()**
  - **mutate()**
  - **summarise()** (along with **group\_by()**)
- Usefully combined with **%>%** operator